# Stroke Prediction

Submissions are evaluated on area under the ROC curve between the predicted probability and the observed target.

**Files**

- train.csv - the training dataset; stroke is the binary target
- test.csv - the test dataset; your objective is to predict the probability of positive stroke
- sample_submission.csv - a sample submission file in the correct format

**Submission File**

For each id in the test set, you must predict the probability for the target variable stroke. The file should contain a header and have the following format:

id, stroke

15304, 0.23

15305, 0.55

15306, 0.98

Etc.